

Logistic regression

Dataset:

German Credit

Objective

Estimate default probabilities using logistic regression

1. Load Libraries and data

In [2]:

```
#Load Libraries
import pandas as pd
import numpy as np
import matplotlib as mpl
import matplotlib.pyplot as plt
from sklearn.cross_validation import train_test_split
from sklearn import linear_model
import statsmodels.api as sm
from sklearn import metrics
from sklearn import datasets
import seaborn as sn
%matplotlib inline
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\cross_validation.py:41: DeprecationWarning: This module was deprecated in version 0.18 in favor of the model_selection module into which all the refactored classes and functions are moved. Also note that the interface of the new CV iterators are different from that of this module. This module will be removed in 0.20.
  "This module will be removed in 0.20.", DeprecationWarning)
```

In [3]:

```
#Load data
credit_df = pd.read_excel('GermanCredit.xlsx')
```

In [4]:

```
#Print header of the file  
credit_df.head()
```

Out[4]:

	Creditability	CreditAmount	DurationOfCreditInMonths
0	1	1049	18
1	1	2799	9
2	1	841	12
3	1	2122	12
4	1	2171	12

2. Check how many records do we have

In [5]:

```
credit_df.shape
```

Out[5]:

```
(1000, 3)
```


3. Plot Histogram for column 'CreditAmount'

In [6]:

```
plt.hist(credit_df['CreditAmount'], 50)
```

Out[6]:

```
(array([ 26.,  87., 124., 121.,  87.,  84.,  63.,  54.,  43.,  49.,  31.,
        24.,  19.,  16.,   7.,  16.,  22.,  14.,  11.,  17.,  13.,   7.,
         7.,   4.,   7.,   5.,   2.,   5.,   3.,   4.,   1.,   4.,   4.,
         2.,   3.,   1.,   0.,   2.,   3.,   2.,   1.,   0.,   3.,   1.,
         0.,   0.,   0.,   0.,   0.,   1.]),
 array([ 250. ,  613.48,  976.96, 1340.44, 1703.92, 2067.4 ,
        2430.88, 2794.36, 3157.84, 3521.32, 3884.8 , 4248.28,
        4611.76, 4975.24, 5338.72, 5702.2 , 6065.68, 6429.16,
        6792.64, 7156.12, 7519.6 , 7883.08, 8246.56, 8610.04,
        8973.52, 9337. , 9700.48, 10063.96, 10427.44, 10790.92,
        11154.4 , 11517.88, 11881.36, 12244.84, 12608.32, 12971.8 ,
        13335.28, 13698.76, 14062.24, 14425.72, 14789.2 , 15152.68,
        15516.16, 15879.64, 16243.12, 16606.6 , 16970.08, 17333.56,
        17697.04, 18060.52, 18424.  ]),
 <a list of 50 Patch objects>)
```



In [7]:

```
amountIntervalsPoints = np.array([0, 500, 1000, 1500, 2000, 2500, 5000, 7500, 10000, 15000, 20000])
amountIntervals = [(amountIntervalsPoints[i] + int(i != 0), amountIntervalsPoints[i + 1]) for i in range(len(amountIntervalsPoints) - 1)]
```

Out[7]:

```
[(0, 500),
 (501, 1000),
 (1001, 1500),
 (1501, 2000),
 (2001, 2500),
 (2501, 5000),
 (5001, 7500),
 (7501, 10000),
 (10001, 15000),
 (15001, 20000)]
```

In [8]:

```
amountIntervalsDf = pd.DataFrame(amountIntervals, columns = ['intervalLeftSide', 'intervalRightSide'])
amountIntervalsDf
```

Out[8]:

	intervalLeftSide	intervalRightSide
0	0	500
1	501	1000
2	1001	1500
3	1501	2000
4	2001	2500
5	2501	5000
6	5001	7500
7	7501	10000
8	10001	15000
9	15001	20000

In [9]:

```
#Credibility table preparation
Credibility0 = []
Credibility1 = []
for interval in amountIntervals:
    subData = credit_df[credit_df.CreditAmount >= interval[0]]
    subData = subData[subData.CreditAmount <= interval[1]]
    Credibility0.append(sum(subData.Creditability == 0))
    Credibility1.append(sum(subData.Creditability == 1))
```

3. Create creditability dataframe

In [10]:

```
tempDf = pd.DataFrame(np.column_stack([Credibility0, Credibility1]), columns = ['Credibiliity0', 'Credibiliity1'])
```

Out[10]:

	Credibiliity0	Credibiliity1
0	3	15
1	34	64
2	51	139
3	33	93
4	26	79
5	75	200
6	34	68
7	20	26
8	21	14
9	3	2

4. Concatenate the above 2 dataframes and give the total of Credibiliity0 and Credibiliity1

In [11]:

```
compareCreditWorthinessDf = pd.concat([amountIntervalsDf.reset_index(drop=True), tempDf], axis=1)
compareCreditWorthinessDf['total'] = compareCreditWorthinessDf.Credibiliity0 + compareCreditWorthinessDf.Credibiliity1
```

Out[11]:

	intervalLeftSide	intervalRightSide	Credibiliity0	Credibiliity1	total
0	0	500	3	15	18
1	501	1000	34	64	98
2	1001	1500	51	139	190
3	1501	2000	33	93	126
4	2001	2500	26	79	105
5	2501	5000	75	200	275
6	5001	7500	34	68	102
7	7501	10000	20	26	46
8	10001	15000	21	14	35
9	15001	20000	3	2	5

5. Plot Creditworthiness plot for Credibility == 0 and

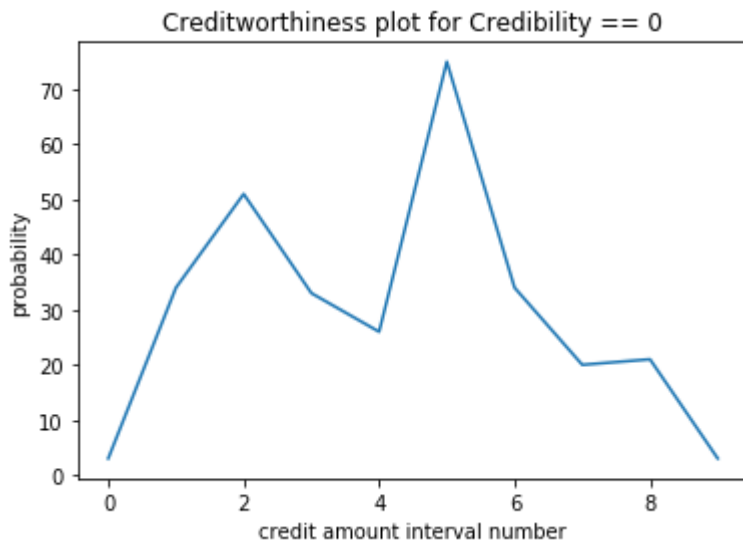
also ==1

In [12]:

```
plt.plot(compareCreditWorthinessDf.Credibiliity0)
plt.xlabel('credit amount interval number')
plt.ylabel('probability')
plt.title("Creditworthiness plot for Credibility == 0")
```

Out[12]:

Text(0.5,1,'Creditworthiness plot for Credibility == 0')




In [13]:

```
plt.plot(compareCreditWorthinessDf.Credibiliity1)
plt.xlabel('credit amount interval number')
plt.ylabel('probability')
plt.title("Creditworthiness plot for Credibility == 1")
```

Out[13]:

Text(0.5,1,'Creditworthiness plot for Credibility == 1')



6. Prepare input data for the model

In [16]:

```
X = np.array(credit_df.CreditAmount)
Y = credit_df.Creditability.astype('category')
```

7. Fit logistic regression model

In [20]:

```
X_train, X_test, y_train, y_test = train_test_split( X, Y, test_size = 0.3, random_state =
logit = sm.Logit( y_train, sm.add_constant( X_train ) )
lg = logit.fit()
lg.summary2()
```

Optimization terminated successfully.
Current function value: 0.598243
Iterations 5

Out[20]:

Model:	Logit	Pseudo R-squared:	0.017
Dependent Variable:	Creditability	AIC:	841.5402
Date:	2018-11-07 21:47	BIC:	850.6424
No. Observations:	700	Log-Likelihood:	-418.77
Df Model:	1	LL-Null:	-425.90
Df Residuals:	698	LLR p-value:	0.00015977
Converged:	1.0000	Scale:	1.0000
No. Iterations:	5.0000		

	Coef.	Std.Err.	z	P> z	[0.025	0.975]
const	1.2287	0.1307	9.4037	0.0000	0.9726	1.4848
x1	-0.0001	0.0000	-3.7764	0.0002	-0.0002	-0.0001

8. Test accuracy calculation

In [22]:

```
def get_predictions( y_test, model ):
    y_pred_df = pd.DataFrame( { 'actual': y_test,
                                "predicted_prob": lg.predict( sm.add_constant( X_test ) ) }
    return y_pred_df

X_test[0:5]
```

Out[22]:

```
array([10974, 1149, 1736, 1414, 2978], dtype=int64)
```

In [23]:

```
y_pred_df = get_predictions(X_test, lg )
y_pred_df['originalCredibility'] = np.array(y_test)
y_pred_df[0:5]
```

Out[23]:

	actual	predicted_prob	originalCredibility
0	10974	0.515795	0
1	1149	0.751504	1
2	1736	0.739680	1
3	1414	0.746211	1
4	2978	0.713492	1

In [24]:

```
y_pred_df['predicted'] = y_pred_df.predicted_prob.map( lambda x: 1 if x > 0.6 else 0)
y_pred_df[0:10]
```

Out[24]:

	actual	predicted_prob	originalCredibility	predicted
0	10974	0.515795	0	0
1	1149	0.751504	1	1
2	1736	0.739680	1	1
3	1414	0.746211	1	1
4	2978	0.713492	1	1
5	2728	0.718888	1	1
6	2859	0.716068	1	1
7	3832	0.694598	1	1
8	727	0.759779	0	1
9	1318	0.748137	1	1


9. Build a confusion matrix

In [26]:

```
def draw_cm( actual, predicted ):
    cm = metrics.confusion_matrix( actual, predicted, [1,0] )
    sn.heatmap(cm, annot=True, fmt='.2f', xticklabels = ["Default", "No Default"] , ytickl
plt.ylabel('True label')
plt.xlabel('Predicted label')
plt.show()
```


In [28]:

```
draw_cm( y_pred_df.originalCredibility, y_pred_df.predicted )
```



In [30]:

```
print( 'Total Accuracy : ', np.round( metrics.accuracy_score( y_test, y_pred_df.predicted ),
```

Total Accuracy : 0.7

10. Predicted Probability distribution Plots for Defaults and Non Defaults

In [29]:

```
sn.distplot( y_pred_df[y_pred_df.originalCredibility == 1]["predicted_prob"], kde=False, co  
sn.distplot( y_pred_df[y_pred_df.originalCredibility == 0]["predicted_prob"], kde=False, co
```

C:\ProgramData\Anaconda3\lib\site-packages\matplotlib\axes_axes.py:6462: UserWarning: The 'normed' kwarg is deprecated, and has been replaced by the 'density' kwarg.

warnings.warn("The 'normed' kwarg is deprecated, and has been "

C:\ProgramData\Anaconda3\lib\site-packages\matplotlib\axes_axes.py:6462: UserWarning: The 'normed' kwarg is deprecated, and has been replaced by the 'density' kwarg.

warnings.warn("The 'normed' kwarg is deprecated, and has been "

Out[29]:

<matplotlib.axes._subplots.AxesSubplot at 0x14e0c2ed128>

